

# GRANTRADAR — PRODUCT REQUIREMENTS DOCUMENT

AI-Powered Grant Intelligence Platform for Research Labs

## Product Overview

**Problem:** Research labs miss critical deadlines because grants are discovered too late. Manual checking is reactive and inefficient. By the time a lab finds a perfect grant, they have only days to apply.

**Solution:** GrantRadar monitors 60+ sources in real-time and alerts labs within minutes of new grant postings. AI predicts when grants will be posted based on historical patterns, giving labs advance notice to prepare applications.

## Core Features (MVP)

<b>1. Real-Time Discovery</b> <ul style="list-style-type: none"><li>Continuous monitoring (15-min intervals)</li><li>NIH/NSF/Grants.gov webhooks + polling</li><li>Adaptive scraping of 50+ foundations</li><li>Alert labs within 5 minutes of posting</li></ul>	<b>2. Predictive Intelligence</b> <ul style="list-style-type: none"><li>Analyze 5+ years of posting patterns</li><li>Predict grant release dates (<math>\pm 3</math> days)</li><li>Alert labs 1-2 weeks before expected</li><li>"Prepare now" mode for draft writing</li></ul>
<b>3. Instant Matching &amp; Alerts</b> <ul style="list-style-type: none"><li>Match new grants within 60 seconds</li><li>Email + SMS + Slack notifications</li><li>Match score with reasoning</li><li>Live dashboard updates (WebSocket)</li></ul>	<b>4. Smart Profile &amp; Learning</b> <ul style="list-style-type: none"><li>5-minute onboarding + ORCID import</li><li>Continuous learning from interactions</li><li>Success prediction scoring</li><li>Competitor tracking (peer lab grants)</li></ul>

## Data Sources & API Integration

### Federal Grant Sources (Primary)

Source	Coverage	API/Method	Update Freq	Historical
NIH RePORTER	\$42B/year, 50k+ active	REST API (free, no auth)	Weekly	1985-present
NSF Award Search	\$8B/year, 11k+ active	REST API (free, no auth)	Daily	1989-present
Grants.gov	All 26 federal agencies	REST API + RSS feeds	Real-time	2002-present
DOD SBIR/STTR	\$4B/year	Web scraping (structured)	Quarterly	2015-present

### Foundation Sources (50+ tracked)

#### Top 20 Priority Foundations:

- Chan Zuckerberg Initiative (\$1B+/year) - API available for some programs, otherwise scrape
- Bill & Melinda Gates Foundation (\$6B+/year) - Structured HTML scraping
- Simons Foundation (\$500M+/year) - RSS feed + web scraping
- Howard Hughes Medical Institute (\$800M+/year) - Structured scraping
- Wellcome Trust (\$1.5B+/year) - REST API available
- Alfred P. Sloan Foundation, Templeton Foundation, Arnold Ventures, Open Philanthropy, Moore Foundation, Kavli Foundation, Schmidt Futures, Carnegie Corporation, Rockefeller Foundation, Ford Foundation (15 more tracked)

## Real-Time Tracking Workarounds

**Challenge:** Not all sources update in real-time. NIH RePORTER updates weekly, some foundations update monthly, DOD SBIR posts quarterly. We need multi-source triangulation to maintain <5 minute alerting.

## Multi-Source Triangulation Strategy

Problem	Workaround Strategy
NIH RePORTER <i>Weekly updates only</i>	Primary: Grants.gov RSS feed catches NIH postings in real-time (NIH posts to both simultaneously) Secondary: Direct scraping of grants.nih.gov/funding-opportunities (HTML changes tracked hourly) Tertiary: Email monitoring - NIH sends listserv emails immediately. Set up dedicated inbox, parse emails via API <b>Result: Detect NIH grants within 15 minutes despite API lag</b>
Foundation Sites	Primary: RSS feeds where available (CZI, Simons, Arnold Ventures have feeds)

Monthly updates, irregular	Secondary: HTML change detection - hash entire opportunities page, re-scrape when hash changes (every 30 min) Tertiary: Social media monitoring - foundations announce on Twitter/LinkedIn before website. Use Twitter API v2 to monitor @CZI_Science, @gatesfoundation Quaternary: Newsletter parsing - subscribe to all foundation newsletters, parse incoming emails for grant announcements <b>Result: Detect foundation grants within 1 hour of announcement</b>
<b>DOD SBIR/STTR</b> <i>Quarterly only</i>	Primary: SBIR.gov has daily updates for all agencies (not just DOD) - use this instead Secondary: Pre-solicitation notices on SAM.gov (beta.sam.gov API) - agencies post intent 30-60 days early Tertiary: Agency-specific portals (Navy: navysbir.com, Air Force: afwerx.com) often post before consolidated site <b>Result: Detect SBIR opportunities 30+ days before official posting via pre-solicitations</b>
<b>European Sources</b> <i>ERC, Horizon, etc.</i>	Primary: Horizon Europe Portal (ec.europa.eu/info/funding-tenders/opportunities/portal) - has RSS feeds Secondary: ERC (erc.europa.eu) updates quarterly but announces via press releases first - scrape press page <b>Result: Track international opportunities for US labs with EU collaborators</b>

## Change Detection Architecture

**HTML Fingerprinting:** For each source URL, compute SHA-256 hash of filtered HTML (remove timestamps, session IDs, ads). Store hash in Redis. Poll every 15-30 min. If hash changes → trigger full scrape + LLM extraction. Avoids expensive scraping when nothing changed.

**Differential Scraping:** When change detected, use LLM to compare old vs new HTML: 'What grants were added or changed?' This avoids re-processing entire catalog. LLM responds: 'New grant: R01-AI-123456, title changed for R21-HL-789012'. Only process deltas.

**Email Parsing Pipeline:** Set up grantradar@yourdomain.com. Subscribe to: NIH Guide (list.nih.gov), NSF funding alerts, all foundation newsletters. Use Gmail API or IMAP to pull emails. LLM extracts structured data: 'From: grants@nih.gov, Subject contains PA-25-123 → parse as new grant announcement'. 90% of foundations announce via email before website.

**Social Media Monitoring:** Twitter API v2: Monitor @NSF, @NIH, @CZI\_Science, @gatesfoundation, @wellcometrust. Set up keyword alerts: 'funding opportunity', 'RFP', 'call for proposals', 'now open'. When match found, extract URL → trigger immediate scrape. Average lead time: 30-120 minutes before RSS/website.

## Redundancy & Failover System

**Multi-Path Verification:** Every grant must be confirmed by 2+ sources before alerting users. Example: NIH R01 appears in (1) Grants.gov RSS, (2) NIH Guide email, (3) grants.nih.gov scrape. If only 1 source shows it → flag for manual review (could be parsing error). If 2+ sources confirm → send alert immediately.

**Source Priority Ranking:** Tier 1 (most reliable): Official APIs (NSF, Grants.gov). Tier 2: RSS feeds, official emails. Tier 3: Web scraping. Tier 4: Social media. Use Tier 1 as ground truth. If Tier 3/4 detects something not in Tier 1/2 within 6 hours → alert human reviewer before sending to users.

**Latency Monitoring:** Track time delta: actual\_post\_time (from agency) vs our\_detection\_time. Goal: median <5 min, 95th percentile <30 min. If detection latency >2 hours for any source → alert on-call engineer. Dashboard shows: 'NIH: 3 min avg, NSF: 8 min avg, CZI: 45 min avg'.

**Backup Scraping:** If primary method fails (API down, rate limit hit, parsing error), automatically fallback to alternative. Example: Grants.gov API down → switch to RSS → if RSS fails → scrape HTML directly. Never miss a grant due to single point of failure.

## Implementation Priority

**Week 1 MVP:** Grants.gov RSS (covers 80% of federal), NSF API, basic NIH scraping. Get to 80% coverage fast.

**Week 2-3:** Add email parsing (NIH Guide, foundation newsletters), HTML change detection, top 10 foundation scrapers.

**Week 4+:** Social media monitoring, pre-solicitation tracking (SAM.gov), European sources, multi-path verification.

## API Implementation Details

**NIH RePORTER API:** Endpoint: <https://api.reporter.nih.gov/v2/projects/search> | Rate limit: 3 req/sec | Fields: project\_num, title, abstract, pi\_name, org\_name, award\_amount, fiscal\_year, project\_start, project\_end | Historical: 40+ years, 3M+ grants

**NSF API:** Endpoint: <https://www.research.gov/awardapi-service/v1/awards.json> | Rate limit: None (pagination: 25/page) | Fields: id, title, abstract, pi, institution, amount, start\_date, expiration\_date | Historical: 1989-present, 500k+ grants

**Grants.gov API:** Endpoint: <https://api.grants.gov/v1/api/search2> | Webhooks: Available via RSS | Fields: opportunity\_id, title, description, eligibility, deadline, funding\_amount, agency | Real-time: New postings appear within minutes

**Wellcome Trust API:** Endpoint: <https://api.wellcome.org/grants> | Auth: API key required (free) | Fields: grant\_id, title, description, funding\_orgs, amount\_awarded, award\_date | Historical: 2005-present

## Grant Prediction Methodology

### Historical Data Collection (Backtesting)

**Data Requirements:** Minimum 5 years of posting history per grant program. Collect: posting date, deadline, funding cycle, fiscal year, agency schedule changes, congressional budget timing.

**Step 1 - Historical Scraping:** Pull all NIH funding opportunity announcements (FOAs) 2019-2024 from RePORTER. Extract: FOA number, posting date, parent announcement. Build database of ~2,000 recurring programs with 5-10 postings each.

**Step 2 - Pattern Recognition:** Group by grant mechanism (R01, R21, K99, etc.). Identify posting patterns: (a) Fixed schedule - R01 typically 3x/year: Jan, May, Sept; (b) Fiscal calendar - many agencies post in Q4 (Jul-Sep); (c) Day-of-week patterns - Wednesdays most common; (d) Lead time - average 60-90 days between posting and deadline.

**Step 3 - LLM Analysis:** Feed historical data to Claude with prompt: 'Analyze posting patterns. Identify: (1) Regular vs irregular programs, (2) Seasonal patterns, (3) Dependencies (budget cycles, leadership changes), (4) Anomalies (COVID delays, policy shifts).' Output: Structured JSON with confidence scores.

**Step 4 - Prediction Model:** Train simple time-series model per grant program. Features: days\_since\_last\_post, month, quarter, fiscal\_year\_day, weekday, historical\_variance. Target: days\_until\_next\_post. Model: LightGBM or simple moving average. Accuracy target: 85% within ±3 days for regular programs.

## Real-Time Prediction Updates

**Continuous Learning:** Every time a predicted grant is posted, log: predicted\_date, actual\_date, error\_days. Retrain model monthly with new data. Adjust confidence scores based on recent accuracy.

**Confidence Tiers:** High (>85% accuracy): Programs with 5+ years consistent pattern. Alert 14 days before. Medium (70-85%): Some variance. Alert 7 days before. Low (<70%): Irregular programs. No prediction alerts, only react when posted.

**Alert Strategy:** 14 days before: 'NIH R01 expected Jan 15 (±3 days, 90% confidence). Start drafting now.' When posted: 'NIH R01 just posted! We predicted this 12 days ago. Your draft should be ready.'

## Validation & Testing

**Backtesting Approach:** Train on 2019-2022 data. Test predictions on 2023-2024 actual postings. Measure: mean absolute error (MAE), accuracy within ±3/7/14 days, false positive rate (predicted but never posted).

**Success Criteria:** 85% of high-confidence predictions within ±3 days. 70% of medium-confidence within ±7 days. <5% false positives. Users find predictions valuable (measured by 'prepare now' alert engagement >40%).

**MVP Validation:** Start with top 20 most predictable programs (NIH R01, NSF CAREER, etc.). Manual validation for first 3 months. Collect user feedback: 'Was this prediction helpful?' Expand to more programs as accuracy improves.

## Agentic Architecture & Backend Workflows

The platform runs on 5 autonomous AI agents that operate 24/7 without human intervention. Each agent has specific responsibilities and communicates through a shared event bus.

Agent	Responsibilities & Workflows
<b>Discovery Agent</b> <i>Real-time monitoring</i>	<ul style="list-style-type: none"><li>Continuous Polling: Check NIH/NSF/Grants.gov every 15 minutes. Compare checksums to detect new postings instantly.</li><li>Webhook Listeners: Subscribe to RSS feeds, API webhooks where available. Receive push notifications for new grants (NIH uses this).</li><li>Adaptive Scraping: LLM-powered HTML parsing every 30 minutes for foundations. Auto-adjusts if page structure changes. Retries with variations.</li><li>Speed Target: New grant detected → processed → matched → alert sent in &lt;5 minutes end-to-end.</li></ul>
<b>Prediction Agent</b> <i>Forecasts grant postings</i>	<ul style="list-style-type: none"><li>Historical Analysis: Analyze 5+ years of posting dates for every grant program (e.g., NIH R01 typically posts 3rd Wednesday of Jan/May/Sept).</li><li>Pattern Recognition: LLM identifies patterns - day of week, month, quarter, federal fiscal calendar, agency-specific schedules.</li><li>Prediction Window: Forecast posting date ±3 days with 85% confidence. Alert labs 7-14 days before predicted date.</li><li>Preparation Mode: Send "Start drafting now" alerts with previous year's RFP, eligibility criteria, and suggested outline.</li></ul>
<b>Curation Agent</b> <i>Validates &amp; enriches data</i>	<ul style="list-style-type: none"><li>Quality Check: LLM scores each grant 0-100 for completeness (has deadline? amount? description?). Flags low-quality for manual review.</li><li>Data Enrichment: If deadline missing, agent visits grant URL and extracts it. Fills eligibility, focus areas, requirements automatically.</li><li>Deduplication: Embeddings + LLM to find same grant from multiple sources. Merges entries, keeps most complete version.</li><li>Auto-Categorization: Assigns research fields, grant types, career stages, keywords. Creates vector embeddings for semantic search.</li></ul>
<b>Matching Agent</b> <i>Connects grants to labs</i>	<ul style="list-style-type: none"><li>Profile Building: Analyzes lab's publications (PubMed API), past grants (NIH RePORTER), and onboarding answers. Creates semantic profile embedding.</li><li>Semantic Matching: Compares grant embedding to lab profile (cosine similarity). LLM evaluates fit: research alignment, eligibility, strategic value.</li><li>Success Prediction: Finds similar labs that applied to similar grants (historical data). Calculates win probability.</li></ul>

	<p>probability based on outcomes.</p> <ul style="list-style-type: none"> <li>Strategic Insights: LLM generates personalized advice: why apply, what to emphasize, potential collaborators, timeline recommendations.</li> </ul>
<b>Delivery Agent</b> <i>Instant multi-channel alerts</i>	<ul style="list-style-type: none"> <li>Instant Notifications: Push alerts via email, SMS (Twilio), Slack webhook within 60 seconds of matching. Priority queue: &gt;90% match = immediate.</li> <li>Smart Routing: LLM decides urgency level. Critical (&gt;95% match, &lt;14 days deadline) = SMS+Email+Slack. High (&gt;85%) = Email+Slack. Medium = Email only.</li> <li>Live Dashboard: WebSocket connection pushes updates to browser. New grants appear with animation. Real-time counter: "3 new grants today".</li> <li>Personalized Content: Each alert includes match score, key reasons, deadline countdown, strategic advice. Never generic.</li> </ul>
<b>Orchestrator Agent</b> <i>Real-time coordination</i>	<ul style="list-style-type: none"> <li>Event Pipeline: Grant discovered → validate (30s) → match all labs (60s) → send alerts (30s). Total latency &lt;2 minutes for 100 labs.</li> <li>Priority Queue: Critical grants jump queue. &gt;95% match + &lt;30 days = highest priority. Background tasks (re-indexing, analytics) run during low traffic.</li> <li>Health Monitoring: Real-time dashboards for agent latency, error rates, queue depth. Auto-scale workers when queue &gt;100 jobs. Alert on-call if any agent down &gt;5min.</li> <li>Circuit Breakers: If NIH API fails, switch to backup scraper. If Claude API slow (&gt;10s), route to GPT-4. Graceful degradation always.</li> </ul>

## Backend Implementation Details

**Real-Time Event Stream (Redis Streams):** Pub/sub architecture with consumer groups. Events: grant\_discovered, grant\_validated, match\_computed, alert\_sent. Each event timestamped for latency tracking. Target: end-to-end <2 min.

**Priority Job Queue (Celery + RabbitMQ):** 3 priority levels: critical (>90% match), high (new grants), normal (re-indexing). Workers auto-scale 3-20 based on queue depth. Route critical jobs to fastest workers.

**Vector Database (Pinecone):** Real-time upserts on new grants (<100ms). Streaming similarity search across 10k+ lab profiles in <200ms. Namespace per user + global index for batch queries.

**WebSocket Server (Socket.io):** Persistent connections to dashboard. Push updates on: new grants, deadline reminders, prediction alerts. Reconnection with exponential backoff. Redis pub/sub for horizontal scaling.

**LLM Orchestration:** Claude Sonnet 4 for matching/insights (quality). GPT-4o-mini for extraction (speed). Groq Llama for categorization (throughput). Latency SLA: 95th percentile <3s. Automatic failover between providers.

## User Flows

**Onboarding (5 min):** Sign up → Answer 8 questions → Import ORCID → Agents build profile in background → First alerts within 2 hours (or immediately if perfect match found)

**Daily Engagement:** Receive instant alerts (SMS/Email/Slack) → Open dashboard (live updates) → Review grant details → Save/dismiss → Dashboard shows "3 new, 5 saved, 2 deadlines this week"

Success Metrics	Tech Stack
<ul style="list-style-type: none"> <li>Alert Latency: &lt;5 min discovery → alert</li> <li>Daily Active Users (DAU) &gt;40%</li> <li>Alert open rate &gt;70% (SMS/Email)</li> <li>Match accuracy &gt;85% (user saves)</li> <li>Prediction accuracy &gt;80% (<math>\pm 3</math> days)</li> <li>Churn &lt;8%/month</li> <li>NPS &gt;55</li> </ul>	<ul style="list-style-type: none"> <li>Frontend: React + WebSocket (live updates)</li> <li>Backend: Python (FastAPI) + Socket.io</li> <li>Database: PostgreSQL + Pinecone (vector)</li> <li>Queue: Celery + RabbitMQ (priority queue)</li> <li>Events: Redis Streams (real-time)</li> <li>AI: Claude/GPT/Groq with failover</li> <li>Scraping: Playwright + Webhooks</li> <li>Notifications: Twilio (SMS) + SendGrid</li> <li>Hosting: Vercel (FE) + Railway (BE)</li> </ul>

## MVP Timeline (3 Weeks)

Week	Deliverables
<b>Week 1</b>	Real-time infrastructure: Redis Streams event bus. Celery + RabbitMQ priority queue. PostgreSQL + basic schema. Build polling scrapers for NIH/NSF APIs (15-min intervals). Test end-to-end latency.
<b>Week 2</b>	Matching + alerts: TF-IDF matching engine. Email (SendGrid) + SMS (Twilio) integration. Simple onboarding form. Manual profile creation for 5 beta users. Test: grant posted → alert sent in <5 min.

<b>Week 3</b>	Live dashboard: React + WebSocket for real-time updates. Grant cards with match scores. Prediction Agent v1 (basic historical pattern analysis). Landing page + Stripe. Launch with 10 paying users at \$200/mo. Validate <5 min alerts.
---------------	------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

## Out of Scope (V1)

- Application writing assistance or grant proposal tools
- Collaboration features (team workspaces, shared pipelines)
- Mobile app (responsive web only)
- Integrations with institutional systems (InfoEd, Cayuse)
- Advanced analytics dashboard (heat maps, trend analysis)